

Kathleen Bartin

AI TRANSFORMATION · ADOPTION & ENABLEMENT · VALUE REALIZATION

Millbrook, NY (New York metro) · 314-799-1248
kathleen@prodigalparadigm.com · prodigalparadigm.com
linkedin.com/in/kathleen-bartin · github.com/prodigalparadigm
Available immediately · New York City · hybrid, remote, or relocation

77.1% of the organization actively engaged — opt-in, no mandate	\$194.5K–\$387.5K documented value delivered on a single engagement	\$156K–\$421K estimated recurring savings a year, every line owned	12 public repositories; four suites run CI on every push
--	--	---	---

SUMMARY

I lead enterprise AI adoption and build the systems that make it hold. At a regulated life insurer this summer I took organization-wide engagement from essentially zero to 77.1% in three months — opt-in, no mandate, part-time — against the 50% leadership said at close it would have been very pleased to reach. I owned the engagement end to end: adoption strategy, governance advice the Chief Legal Officer acted on, the client's first production AI agent, and a priced value case. It ended in succession, not dependency: the client's senior AI engineer compiled my enablement library into version-controlled skills and now runs the internal build practice. This is the second general-purpose technology I have helped organizations adopt.

EXPERIENCE

AI Enablement Consultant · Ceres Life Insurance Company May 2026 – August 2026
Contract via Prodigal Paradigm LLC · regulated life insurance · New York metro · CITO-sponsored · won, scoped, priced and delivered solo

- ♦ **Adoption program.** Designed and ran enterprise-wide adoption — stakeholder discovery, process mapping, role-based training, cohort rollout, office hours, playbooks, a champions network — with weekly value reporting to the sponsor; trained and advised the full C-suite. Group walkthroughs proved the highest-leverage tactic, reaching shift-based and remote staff; the data showed that a manager's own non-engagement blocks an entire reporting line.
- ♦ **Governance with Legal.** Advised the Chief Legal Officer on model-level zero-data-retention differences; the client acted on it, removing a model from the enterprise picker and moving to the platform tier that enforces retention. Wrote the HR playbook drawing the drafting-versus-screening line under automated-employment-decision rules.
- ♦ **First production agent.** Shipped the client's first production AI agent — Slack help-desk triage on a 30-minute cycle — under a one-page governance package: data classification, human-in-the-loop controls, logging, responsible-use boundaries and a kill switch. Ran it observe-first, then validated it in supervised live runs.
- ♦ **Builds.** A fully local call-center QA scoring system (Whisper, diarization, semantic and paralinguistic tracks; \$40,000–\$80,000 vendor-equivalent), built with no data-science team, then rebuilt independently by the operations manager; a three-agent QA harness — tester, evaluator, orchestrator — testing software on synthetic data, handed to the client's dev team on request; and scheduled automation (live leadership adoption dashboard, daily comms triage, weekly tracker reconciliation) through native Microsoft 365 connectors, Teams, Slack, SharePoint and ClickUp.
- ♦ **Operational risk.** Diagnosed a multi-day organization-wide AI-platform outage — agentic workloads exhausting shared rate limits and starving interactive use — and designed the remediation separating agentic from human capacity, with approval gates. Brought outsourced legal contract work back in-house with an AI-assisted workflow.
- ♦ **Value, priced.** Priced the engagement layer by layer from the complete activity log: \$194,500–\$387,500 delivered and \$156,000–\$421,000 a year in estimated recurring savings, every line with stated assumptions and a named owner.
- ♦ **Succession, not handoff.** Left 50+ knowledge documents, 15+ training videos, role-based curricula, shared Claude Projects, packaged platform skills and a 56-resource enablement library with three named internal owners. In the final fortnight the client's senior AI engineer took the library unprompted, compiled it into version-controlled department skills with compliance guardrails, credited the original authorship, and now holds the internal AI build practice.

Principal · Prodigal Paradigm LLC January 2025 – Present
AI adoption consulting, products and education · Millbrook, NY

- ♦ **Public tool line** at prodigalparadigm.com/tools — essays and methods turned into working tools: **Answerable**, an AEO audit of how AI answer engines read a website; **The Supersuit Check**, an AI-posture self-assessment; **A Week with Fred**, a simulator for managing a brilliant agent with total amnesia; **Parrot**, a screening-call copilot suggesting three next questions every 20 seconds (pre-release).
- ♦ **The Recruiter's Field Kit** — open-source decoder that reads an AI-era job description, counts the distinct labor markets inside it, and writes ten-minute screening questions with the real answer and the bluff to listen for. Role-count agreement with my own read rose from 2 of 6 to 8 of 10; its central rule — no automated LinkedIn access, ever — is enforced by CI on every push.
- ♦ **Claude plugins** that package my method: **practice-harvest** finds the repeatable work in your own history and crystallizes it into skills; **mywork-charter** sorts which work stays yours and charts what the freed time is for.
- ♦ **Build-to-order systems** on my own stack: a full-network CRM with correspondence and pipeline, built in a day; memory and retrieval systems; QA harnesses for agents and prompts; house-style skills; and **RAVEN**, an agent that learns taste and hunts dealer inventory.
- ♦ **Community teaching.** Founded **Coffee with Claude** at the Millbrook Public Library (standing since June 2026), packaged as **Coffee with Claude in a Box** so any library can run it; launching **The Doctor Is In**, an AI-help booth, at community events this fall.

Director of Client Relations (part-time) · LaUL Gallery

June 2024 – July 2026

- ◆ Built a \$200,000+ client book from a standing start, selling fine jewelry at \$500–\$25,000+ per transaction: 80% close rate, 50+ high-net-worth accounts, 43% year-over-year transaction growth. Founded client relations; recruited and onboarded my successor.

Foundation Services Specialist · The Chamber Foundation

April 2025 – November 2025

- ◆ Designed an eight-lesson youth-leadership curriculum and wrote approved one-, two- and three-day workforce bootcamp curricula; orchestrated 80+ volunteers across four committees.
- ◆ **Grant writing.** Wrote 20 grant requests, including a six-letter cold donor campaign against a \$35,000 sponsorship shortfall; one letter converted a private charitable trust into a first-time funder of the youth teambuilding day at the FDR Presidential Library — funded almost a year after I left.

Prospect Researcher to Senior Researcher · Washington University in St. Louis

2009 – 2012

- ◆ Exceeded annual reporting quota by 400% on a multibillion-dollar capital campaign; pioneered international prospecting and social-media research methods that later became industry standard.

Earlier: Technology Trainer and Sales, New Horizons / Productivity Point International — sold and delivered enterprise computing training through the first wave of corporate adoption. 2012–2024: full-time parent, with intermittent contract roles.

BUILDS, RESEARCH AND WRITING

Twelve public repositories (github.com/prodigalparadigm, github.com/comelyhunger), built with Claude. The four MIT suites below postdate the Ceres engagement and run offline with no credentials; CI runs pytest, ruff and strict mypy on every push.

- ◆ **servicenow-mcp** — MCP server over ServiceNow ITSM: field allowlists pushed down as sysparm_fields, record caps, 429-aware token-bucket rate limiting, read-only by default at two layers, and correlation-id idempotency so a replayed create never files twice.
- ◆ **stagegate** — runs each agent capability at a declared stage — observe, suggest, act — from a reviewed TOML policy, with a hash-chained audit record per call, redaction before any sink, fail-closed approvals, and a kill switch that degrades to shadow mode.
- ◆ **promptproof** — synthetic conversations across six axes, including multi-turn drift ladders that install a premise before cashing it in; an LLM-judge rubric reports per model which instruction broke, and on which turn.
- ◆ **callscope** — fully local two-speaker call QA: energy VAD with hysteresis, MFCC plus log-FO diarization with a silhouette guard against false splits, and semantic and paralinguistic tracks scored independently.
- ◆ **Lab 01 — Give Claude a Tool** — 30-minute hands-on GTM lab on the tool-use loop, with a pre-class smoke test, captured transcripts that keep the room working through an API outage, instructor notes and a release-cycle maintenance checklist.
- ◆ **Production stack** — ten MCP servers I built and run daily, and hybrid-retrieval RAG (dense plus sparse, reciprocal-rank fusion) over 70,000+ chunks on Qdrant.
- ◆ **Research** — pre-registered longitudinal measures over a 495-conversation, 20,000-turn corpus (Kaggle / Google DeepMind Metacognition track, 2026). NYC AI security hackathon 2026: one prompt, two frontier models — model choice is a safety decision.
- ◆ **Writing** — essays at prodigalparadigm.com, bylined Kathleen Bartin, with Claude — What Counts; Actuaries and LLMs Are Cousins.

METHODS I BRING

- ◆ **The Underwriting Principle:** you would not price a product on unvalidated assumptions; do not automate a workflow on unarticulated knowledge. Articulate, test and validate, then automate.
- ◆ **Adoption maturity ladder:** Equipped, Oriented, Builder, Claude Code-ready, so “engaged” means the same thing every month.
- ◆ **The Rule of Three:** explained three times, it belongs in a shared Project.
- ◆ **Make the client the hero:** every asset usable without me; every engagement designed to end in succession, not dependency.

SKILLS

Adoption and change: AI adoption strategy, change management, champions networks, communities of practice, train-the-trainer, stakeholder discovery, process mapping, use-case prioritization, adoption measurement, value realization and ROI, cross-functional stakeholder management, user onboarding.

Enablement and learning: instructional design, technical curriculum development, hands-on labs and workshops, office hours, AI literacy programs, playbooks and SOPs, executive briefing.

Governance, risk and delivery: responsible AI, AI governance, data classification, human-in-the-loop design, model selection and data-retention posture, agent audit and approval controls, AI evaluation, compliance-aware deployment in regulated industries, AI implementation and forward-deployed client delivery, proof-of-concept (POC) scoping, end-to-end project ownership.

Technical: generative AI (GenAI) and LLM integration; Claude (Enterprise, Cowork, Claude Code, API), ChatGPT, Gemini, prompt engineering, agent design, MCP server development, RAG (Qdrant, ChromaDB), Python, Git and GitHub, Whisper, ffmpeg, Microsoft 365, SharePoint, Teams, Slack, ClickUp.

EDUCATION, CREDENTIALS AND SERVICE

BA, Philosophy — Excelsior University. Coursework: University of Chicago; graduate Actuarial Mathematics, University of Missouri.

Twelve Anthropic Academy certificates (2025–2026); Google AI Essentials (2025). **US Navy**, honorable discharge. Licensed EMT (New York). Board member, Millbrook Business Association; Rotary Exchange Counselor; Millbrook Hook and Ladder.